
SkillForge AI

LLM API Cost & Usage Report

Token Usage | Cost Analysis | Optimization Strategy

June 2026

CONFIDENTIAL | Internal Use Only

1. Executive Summary

Bottom Line: Current Gemini 2.0 Flash setup is already near-optimal. Adding caching + routing simple tasks to Groq llama-3.1-8b cuts billing by 60–70% with zero quality loss.

Metric	Value
Primary LLM	Gemini 2.0 Flash (Google AI)
Fallback LLM	llama-3.3-70b-versatile (Groq)
Total LLM call types	13 distinct functions
Cost per 100 active users/day	~\$0.30 – \$0.45 / day
Monthly cost (100 users)	~\$9 – \$14 / month
Monthly cost (1,000 users)	~\$90 – \$140 / month
Free tier available	YES – Gemini & Groq both have free tiers

2. Current API Configuration

Active API Keys

Key	Provider	Model	Status
GEMINI_API_KEY	Google AI Studio	gemini-2.0-flash	Active – Primary
GROQ_API_KEY	Groq Cloud	llama-3.3-70b-versatile	Active – Fallback
OPENAI_API_KEY	OpenAI	GPT-4o-mini	Configured but unused

Failover Strategy

User Action
-> Gemini 2.0 Flash [Primary]
-> Groq llama-3.3-70b [on failure / 429 quota]
-> Rule-based / Template [on all LLM failure]
-> Response returned [always succeeds]

Key Parameters

Setting	JSON Calls	Text Calls
Max output tokens	2,048	1,024
Temperature	0.7	0.8
Timeout	30 seconds	20 seconds

3. All LLM Calls — Detailed Breakdown

#	Function	API / Model	Input Toks	Output Toks	Trigger
1	Skill Decomposition	Gemini 2.0 Flash	1,500–2,000	500–800	Once per goal
2	Diagnostic Quiz Gen	Gemini 2.0 Flash	1,200–1,800	400–600	Once per goal
3	Session Evaluation	Gemini 2.0 Flash	1,500–2,500	300–500	1–5x/user/day
4	Challenge Generation	Gemini 2.0 Flash	1,200–1,800	600–900	Once/day
5	Progress Report Gen	Gemini 2.0 Flash	1,200–1,600	400–600	On-demand
6	Autonomous Daily Brief	Gemini 2.0 Flash	200–400	100–200	Every 6h
7	Interview Question Gen	Gemini 2.0 Flash	600–900	400–600	Per session
8	Interview Answer Eval	Gemini 2.0 Flash	800–1,200	300–400	5–10x/session
9	Interview Final Report	Gemini 2.0 Flash	1,200–2,000	500–800	Per interview
10	Market Intelligence	Gemini 2.0 Flash	600–900	400–600	On-demand
11	Assessment Gen (JD)	Groq -> Gemini	2,000–4,000	1,000–2,000	Per employee
12	Session Quiz Gen	Gemini 2.0 Flash	1,000–1,500	600–800	Per session
13	Study Notes Gen	Gemini 2.0 Flash	1,000–1,500	600–800	On-demand

Highest Cost Operations

- Assessment Generation (#11): 2,000–4,000 input + 1,000–2,000 output per employee — scales with user count
- Interview Sessions (#7–9): 10,600–18,700 total tokens for a full session
- Tutor Chat (high frequency): 300–1,000 tokens per message, potentially 100+ messages/day/user
- Session Evaluation (#3): Runs 1–5x per active user daily — compounds at scale

4. Token Usage Per Functionality

Per Single User Action

Feature	Input Toks	Output Toks	Total	Cost (Flash)
New Learning Goal	2,700–3,800	900–1,400	3,600–5,200	\$0.0004
Daily Challenge	1,200–1,800	600–900	1,800–2,700	\$0.0001
Submit Response (eval)	1,500–2,500	300–500	1,800–3,000	\$0.0002
Session Quiz	1,000–1,500	600–800	1,600–2,300	\$0.0002
Study Notes	1,000–1,500	600–800	1,600–2,300	\$0.0002
Progress Report	1,200–1,600	400–600	1,600–2,200	\$0.0002
Full Interview Session	8,000–14,000	3,800–6,200	11,800–20,200	\$0.0015
Market Intelligence	600–900	400–600	1,000–1,500	\$0.0001
Tutor Chat (10 msgs)	3,000–10,000	2,000–4,000	5,000–14,000	\$0.0007

Feature	Input Toks	Output Toks	Total	Cost (Flash)
Assessment (20 staff)	40,000–80,000	20,000–40,000	60,000–120,000	\$0.008

Daily Projections by User Volume

Active Users/Day	Daily Tokens	Daily Cost	Monthly Cost
10	~500,000	~\$0.03	~\$0.90
50	~2,500,000	~\$0.15	~\$4.50
100	~5,000,000	~\$0.30	~\$9.00
500	~25,000,000	~\$1.50	~\$45.00
1,000	~50,000,000	~\$3.00	~\$90.00
5,000	~250,000,000	~\$15.00	~\$450.00

Assumptions: avg 2 sessions/day, 5 chat messages, 1 challenge, 1 quiz, 0.1 interviews per user per day

Free Tier Thresholds

Provider	Free Limit	SkillForge Equivalent
Gemini 2.0 Flash	1,000 RPD, 1M TPM	~15–20 active users/day completely free
Groq Cloud	14,400 RPD, 500k TPD	Assessment generation for ~50 employees free/day

5. Cost Calculations

Gemini 2.0 Flash Pricing (June 2026)

Token Type	Price
Input tokens	\$0.075 per 1,000,000 tokens
Output tokens	\$0.30 per 1,000,000 tokens

Scenario A — Startup (100 users, 30% daily active)

Feature	Daily Calls	Input Tokens	Output Tokens	Daily Cost
Skill Decompose	5	8,750	3,250	\$0.0007
Quiz Generate	5	7,500	2,500	\$0.0006
Session Evaluate	30	60,000	12,000	\$0.0082
Challenge Gen	30	45,000	22,500	\$0.0101
Progress Reports	10	14,000	5,000	\$0.0020
Scheduler Brief	10	3,000	1,500	\$0.0007

Feature	Daily Calls	Input Tokens	Output Tokens	Daily Cost
Interview Sessions	5	50,000	23,500	\$0.0108
Market Intel	10	7,500	5,000	\$0.0015
Assessments (JD)	20	60,000	30,000	\$0.0135
Session Quiz	30	37,500	21,000	\$0.0091
Study Notes	20	25,000	14,000	\$0.0060
Tutor Chat	50 msgs	32,500	15,000	\$0.0069
TOTAL/DAY	—	~350,000	~155,000	~\$0.070
TOTAL/MONTH	—	—	—	~\$2.10

Scenario B – Growth (1,000 users): ~\$0.70/day, ~\$21/month

Scenario C – Scale (10,000 users): ~\$4.20/day, ~\$126/month

Assessment Generation Cost

Employees	Questions Each	Gemini Cost	Groq Cost
10	10	\$0.003	\$0.002
50	10	\$0.015	\$0.008
100	10	\$0.030	\$0.016
50	30 (max)	\$0.025	\$0.014

6. LLM Provider Comparison

Pricing Table (Per 1M Tokens — June 2026)

Provider	Model	Input \$/1M	Output \$/1M	Speed	JSON	Free Tier
Google	gemini-2.0-flash	\$0.075	\$0.30	Very Fast	5/5	YES
Google	gemini-1.5-flash-8b	\$0.0375	\$0.15	Fastest	4/5	YES
Google	gemini-2.5-pro	\$1.25	\$10.00	Slow	5/5	Limited
Groq	llama-3.3-70b	\$0.59	\$0.79	Fastest	4/5	YES
Groq	llama-3.1-8b-instant	\$0.05	\$0.08	Fastest	3/5	YES
Groq	gemma2-9b-it	FREE	FREE	Fast	3/5	YES
OpenAI	gpt-4o-mini	\$0.15	\$0.60	Fast	5/5	NO
OpenAI	gpt-4o	\$2.50	\$10.00	Slow	5/5	NO
Anthropic	claude-3-haiku	\$0.25	\$1.25	Fast	5/5	NO
Mistral	mistral-small	\$0.10	\$0.30	Fast	4/5	NO

Provider	Model	Input \$/1M	Output \$/1M	Speed	JSON	Free Tier
Together AI	llama-3.1-8b	\$0.018	\$0.018	Fast	3/5	YES

Best Provider by Use Case

Use Case	Best Provider	Reason
Structured JSON output	Gemini 2.0 Flash	Native JSON mode, very reliable, 1M context
Fast chat responses	Groq llama-3.1-8b	Fastest inference globally, near-free
Complex reasoning	Gemini 2.0 Flash	Best reasoning-to-cost ratio available
Cost-free fallback	Groq free tier	14,400 requests/day completely free
Assessment generation	Groq llama-3.3-70b	High token limit, 4x cheaper than GPT-4o

7. Model Comparison

Feature	gemini-2.0-flash	gemini-1.5-flash-8b	llama-3.1-8b (Groq)	gpt-4o-mini
JSON generation	5/5 BEST	4/5	3/5	5/5
Long context	5/5 (1M ctx)	4/5 (1M ctx)	3/5 (128K)	4/5 (128K)
Educational quality	5/5	4/5	3/5	4/5
Eval accuracy	5/5	4/5	3/5	4/5
Speed	Very Fast	Fastest	Fastest	Fast
Cost	Low	Very Low	Near-Free	Low
Free tier	YES	YES	YES	NO
Overall SkillForge fit	5/5 WINNER	4/5	3/5	4/5

Winner Rankings

Rank	Model	Provider	Monthly (100 users)	Best For
#1 BEST	gemini-2.0-flash	Google	~\$2–9	Everything: JSON, eval, generation
#2	llama-3.1-8b-instant	Groq	~FREE	Simple tasks, chat, scoring
#3	gemini-1.5-flash-8b	Google	~\$1–4	Simple tasks, scheduler briefs
#4	gpt-4o-mini	OpenAI	~\$4–18	JSON if Gemini unavailable
#5	llama-3.3-70b	Groq	~\$3–12	Complex assessment generation

8. Best Strategy — Low Cost + High Performance

Tiered Model Strategy (Recommended)

TIER 1 — Complex / High Quality → Gemini 2.0 Flash (\$0.075/\$0.30 per 1M)

- Skill decomposition
- Assessment generation from JD
- Interview question generation & evaluation
- Open-ended session response evaluation
- Progress report narrative

TIER 2 — Medium Tasks → Groq llama-3.1-8b or Gemini (\$0.05/\$0.08 per 1M)

- Session quiz generation
- Study notes
- Challenge generation
- Market intelligence

TIER 3 — Simple / High Frequency → Groq llama-3.1-8b (Near-FREE)

- Tutor chat responses
- Motivational scheduler briefs
- Simple answer scoring

Cost Savings Opportunities

Optimization	Est. Savings	Effort
Cache assessment questions per JD hash	40–60% on assessment costs	Low
Use llama-3.1-8b for tutor chat	80% on tutor chat costs	Low
Use llama-3.1-8b for scheduler briefs	95% on brief costs	Low
Cache market intelligence (24h TTL)	70% on market intel costs	Low
Reduce tutor chat history 5 → 3 turns	20% on chat token count	Low
Cache quiz templates per skill (7d TTL)	50% on quiz costs	Medium

Projected Savings with Optimizations

Scenario	Current Cost/Month	Optimized Cost/Month	Savings
100 users	~\$9	~\$3	67%
1,000 users	~\$90	~\$28	69%
5,000 users	~\$450	~\$135	70%

9. Recommended Architecture

Immediate (No Code Changes Required)

- Switch tutor chat to Groq llama-3.1-8b — same quality for short responses, near-free
- Switch scheduler briefs to Groq llama-3.1-8b — 2-sentence outputs do not need 70B model
- Enable Gemini free tier monitoring — stay under 1,000 req/day for low-traffic periods

Short Term (1–2 Weeks) — Model Routing Code

```
// server/services/GeminiService.js – Add model routing
const MODEL_TIERS = {
  complex: "gemini-2.0-flash", // Tier 1: JSON, eval, reports
  medium: "gemini-1.5-flash-8b", // Tier 2: quiz, notes, challenges
  simple: "groq:llama-3.1-8b-instant", // Tier 3: chat, briefs, scoring
};
```

Add in-memory cache for:

- Assessment questions per JD hash (TTL: 7 days)
- Market intelligence per domain (TTL: 24 hours)
- Quiz templates per skill (TTL: 24 hours)

Long Term (1 Month+)

- Upgrade to Gemini 2.5 Flash when stable — same price, better reasoning
- Implement token budgets per user per day to cap runaway costs
- Add usage analytics dashboard using GeminiService.callCount + groqCallCount
- Per-company token tracking — store usage in companies record, expose on Super Admin dashboard

10. Immediate Actions

Priority 1 — Do Now (Zero Cost, Instant Savings)

Action	File	Impact
Switch AutonomousScheduler to Groq llama-3.1-8b	server/agent/AutonomousScheduler.js	~95% savings on brief costs
Switch tutor chat to Groq llama-3.1-8b	server/routes/tutor.js	~80% savings on chat costs
Confirm GEMINI_MODEL=gemini-2.0-flash in .env	.env	Already set — verify
Verify Groq free tier	console.groq.com	14,400 free req/day
Verify Gemini free tier	aistudio.google.com	1,000 free req/day

Priority 2 — This Week

Action	Impact
Add LRU cache for market intelligence (TTL 24h)	70% cost reduction for market intel
Add JD hash caching for assessment generation	40–60% reduction in assessment costs

Action	Impact
Reduce tutor chat history from 5 to 3 turns	20% token reduction per message
Log callCount + groqCallCount on /api/health	Usage visibility

Priority 3 — This Month

Action	Impact
Per-company token usage tracking in companies record	Billing visibility per tenant
Usage stats on Super Admin dashboard	Real-time cost monitoring
Daily token budget per company (by plan tier)	Cost cap per company
Evaluate gemini-1.5-flash-8b for medium tasks	50% cheaper than gemini-2.0-flash

Summary — Best Choice for SkillForge AI

Criteria	Winner	Reason
Lowest cost	Groq llama-3.1-8b	Near-free, 14,400 req/day free tier
Best performance	Gemini 2.0 Flash	Best JSON, reasoning, 1M context
Best balance (RECOMMENDED)	Gemini 2.0 Flash + Groq	Already implemented — keep this
Best free option	Gemini free + Groq free	Up to ~30 active users/day free
Best for 1K+ users scale	Gemini 2.0 Flash	\$90/month at 1,000 users — very low
Best for assessment gen	Groq llama-3.3-70b	High token limit, 4x cheaper
Best for chat / tutor	Groq llama-3.1-8b-instant	Fastest, cheapest, good quality

Current setup is already excellent. Switch tutor chat + scheduler briefs to Groq llama-3.1-8b and add result caching to achieve 60–70% billing reduction with ZERO quality loss.

Report generated for SkillForge AI — June 2026
Pricing subject to change — verify at: aistudio.google.com | console.groq.com | platform.openai.com/pricing